

Design Report(Group 14)

Intro

For our project 3 we implemented a B+ tree index to further the Buffer Manager program. Our B+ tree uses a search key of type integer to traverse the tree to find corresponding record id pairs. This allows us to efficiently traverse the database, since it is self-balancing, in $O(\log n)$ for search, insert, and delete.

Traversing the B+ tree

The advantage of using a B+ tree is we can traverse through the tree much like a linked list to find a page/key pair. The nodes contain a pointer to the next node and the key value. Using these nodes, we can traverse the tree in an ordered manner. Since the B+ tree is ordered by left child nodes being less than the parent, right child nodes being greater than the parent, and the middle in-between the two values, we can use this information to efficiently traverse the tree.

Design Choices

We implemented the B+ tree by consistently calling InsertEntry() to insert any new pages. From this point we determine where the entry should go based on the depth of the tree and the balance.

Unpin Page

We decided to unpin the pages as soon as possible. This allows us to keep the buffer pool from filling up. This allows us to speed up the program since the I/O would be a bottle neck, and we are now replacing that with CPU cycles, which are much faster.

Efficiency

Search: searching the tree will result in a I/O cost of $2 \times \text{height}$, since we need to read in and write out through the pages.

Insert: If the root node empty, inserting will be a simple I/O cost of 2, since we are writing directly to the root node of the tree. If it is not empty, we recursively added the new nodes.

Scanning: Scanning the tree will have a worst-case scenario of $O(\log n)$ if we need to read every leaf node, or the same as Search.

Splitting: Splitting results in a I/O cost of 4. This is due to the need of reading in and writing both the current node and the new node we are creating in the split.

Extra Tests

Test 4:

We create an empty B+ tree and run tests to make sure it scans properly. Each call of intScan should throw and catch a NoSuchKeyFoundException and return 0.

Test 5:

Create a B+ tree with tuples valued 0 to 100000 in a random order. Tests if the program will run when tasked with an excessive amount. Might run out of buffer space, if unpin page doesn't work properly.